

Programmierparadigmen

Prof. Dr. Ralf Reussner

Topic 5.1

Parallel Programming Fundamentals

DEPENDABILITY OF SOFTWARE-INTENSIVE SYSTEMS
INSTITUTE OF INFORMATION SECURITY AND DEPENDABILITY, FACULTY OF INFORMATICS

dsis.kastel.kit.edu



Overview on Today's Lecture

- Content
 - Motivation for parallel programming
 - Fundamentals of parallel programming
 - Exemplary parallel algorithms
- Learning goals: participants –
 - understand why parallel programming becomes increasingly relevant
 - understand basic concepts and approaches for parallel programming
 - know the basic terminology in parallel programming
 - can estimate the parallelization potential for given programs
- Parallelization topics in the next weeks
 - C basics & parallel programming with MPI
 - Parallel programming in Java
 - Design by Contract

Computer Performance Development



Atari
Mega ST 2
8 MHz



Umax
Pulsar
150 MHz



Samsung
Omnia II
800 MHz



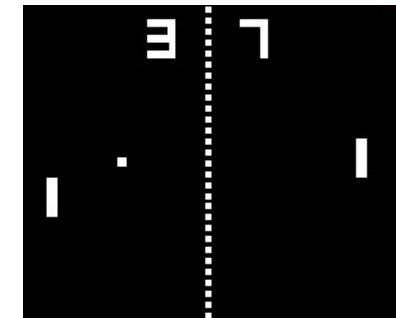
Samsung
Galaxy S25
4470 MHz



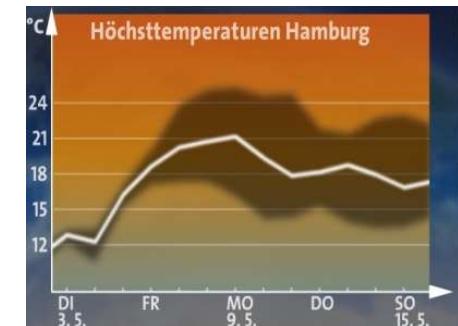
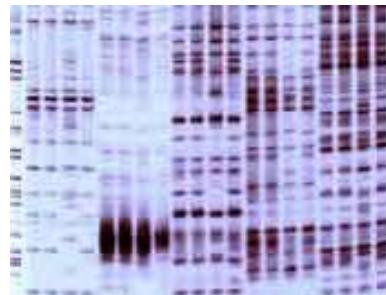
Motivation

- Application of Murphy's Law on computer performance:

Every computer is too slow 😊

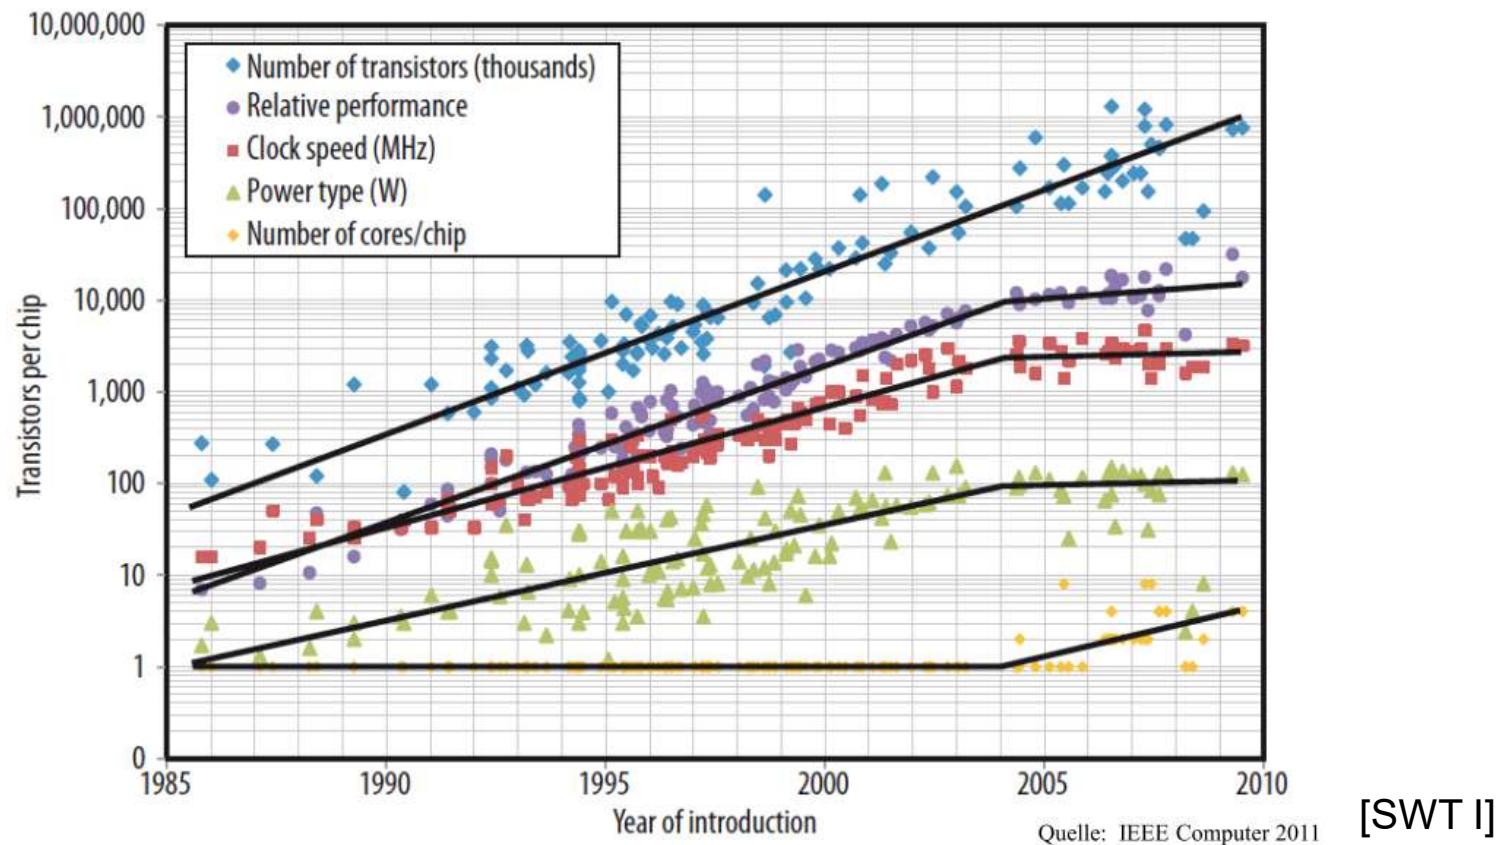


- Why does that bother us?
 - Computer graphics:
Medicine, movies, games, ...
 - Genome analysis
 - Weather forecast
 - etc.



Moore's Law (1)

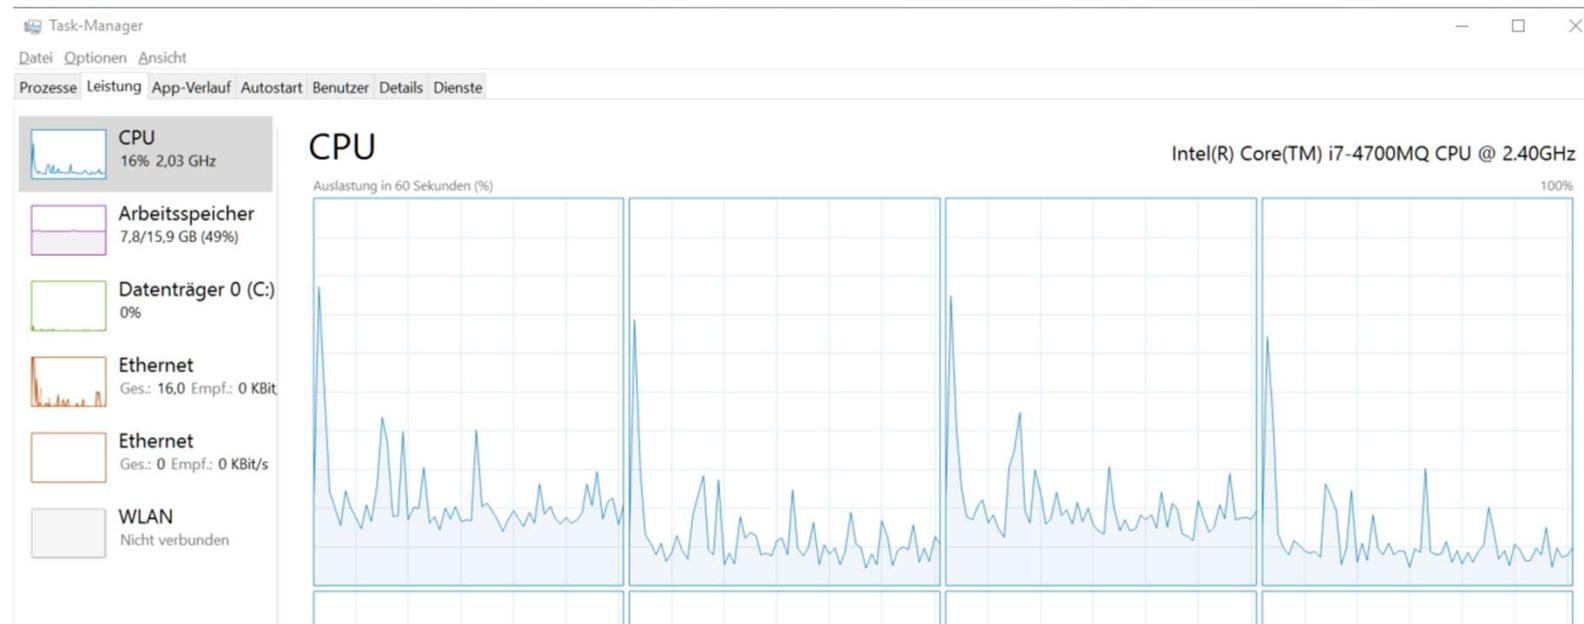
- Moore's law states that the number of transistors in integrated circuits doubles every two years



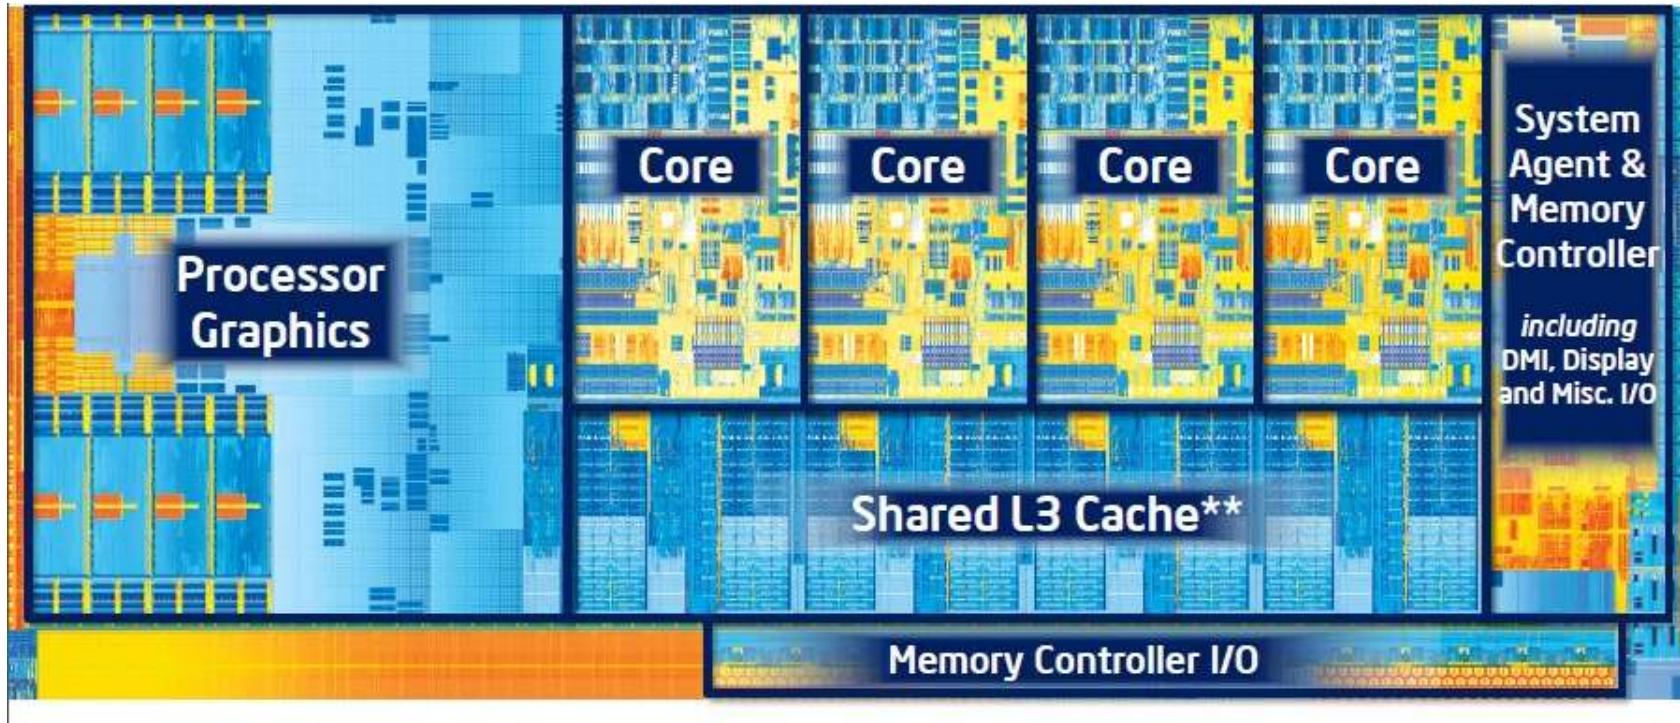
- Meanwhile, Moore's law ceases to be valid

Moore's Law (2)

- Often not only applied on complexity, but also on –
 - performance
 - or even on clock rates
- The clock rates stagnate since several years at about 3 – 4 GHz
- How can performance be further increased?
 - e.g. with **parallel processing** on **multi-core processors** or **distributed systems**



Intel Multicore Architectures



Ivy-Bridge Architecture [Intel]

- 2012 (Ivy-Bridge Desktop): up to 6 Cores
- 2018 (Skylake-X Desktop): up to 18 Cores
- 2025 (Bartlett-Lake Desktop): up to 24 Cores
- Future: ?

Consequences (1)

- Performance can be increased by parallelization
- Main products of processor manufacturers are multi-core processors
 - Servers ship with multi-core processors since 2005
 - Laptops ship with dual-core processors since 2006
 - Even smartphones use multi-core processors since several years
- **Parallel processing becomes the rule**
- Fundamental transition from sequential to parallel processing
 - Dramatic impacts on applications, research and education
 - Buzzwords: Grid and Cloud Computing, Cluster, Supercomputer
- **Computer scientists have to learn how to master „parallel applications“**

Consequences (2)

- Software producers (and the computer science education) have to adapt to parallel processing techniques to stay competitive
 - Migration of existing applications
 - Development of new, parallel applications
- More details on parallel programming in advanced lectures
 - Parallelrechner und Parallelprogrammierung
 - Parallele Algorithmen
 - Effizientes paralleles C++
 - etc.

Parallel Processing

- No programming paradigm in the proper sense
 - Can be implemented in different paradigms and languages
- The following lectures demonstrate basic concepts exemplarily in –
 - C: Message Passing Interface (MPI)
 - Java: Advanced concepts and actors
- Modern processors support parallel execution on different levels
 - Bit level (e.g. two 32-bit additions in a 64-bit architecture)
 - Instruction level (e.g. pipelining)
 - Data level (e.g. loops)
 - Processing level (common in modern systems)

Parallelism vs. Concurrency

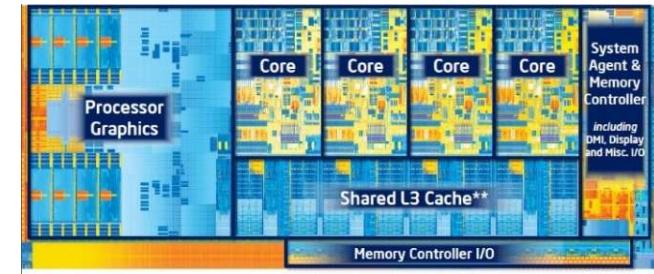
- In the context of parallel computing, the terms **parallelism** and **concurrency** are commonly used
- We follow the definition of Oracle for these terms [Oracle]

- **Parallelism:** At least two threads are *executed simultaneously*
- **Concurrency:** At least two threads *make progress*

- Concurrency is a generalized form of parallelism
 - Includes time-slicing as a form of virtual parallelism, which computers support for many years

Parallel Programming Approaches

- Two different parallel programming approaches
- For computer systems with **shared memory**
 - Each processor can access each memory cell
 - Processors communicate by manipulating shares memory cells
 - e.g. on multi-core processors
- For computer systems with **distributed memory**
 - Each processor has its own memory, which is only accessible by that processor
 - Processors communicate via messages (**message passing**)
 - e.g. on computer clusters



[Intel]



[NASA]

Shared memory (1)

Process

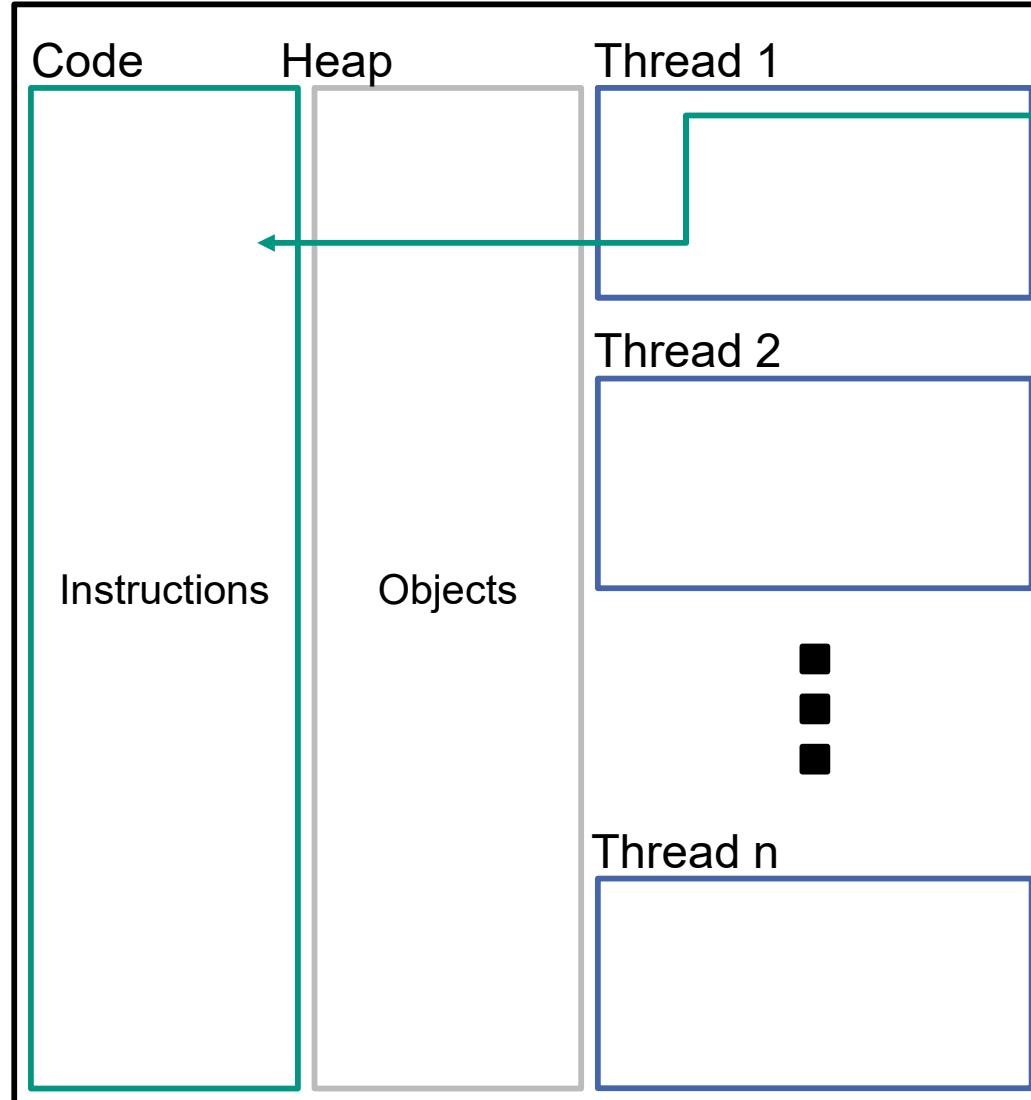
- Generated by **operating system**
- Contains information about **program resources** and **execution status**, e.g.
 - Code segment (program instructions)
 - Data segment (for global variables, stack, heap)
 - At least 1 control thread
- CPU context switch between processes slow

Control thread

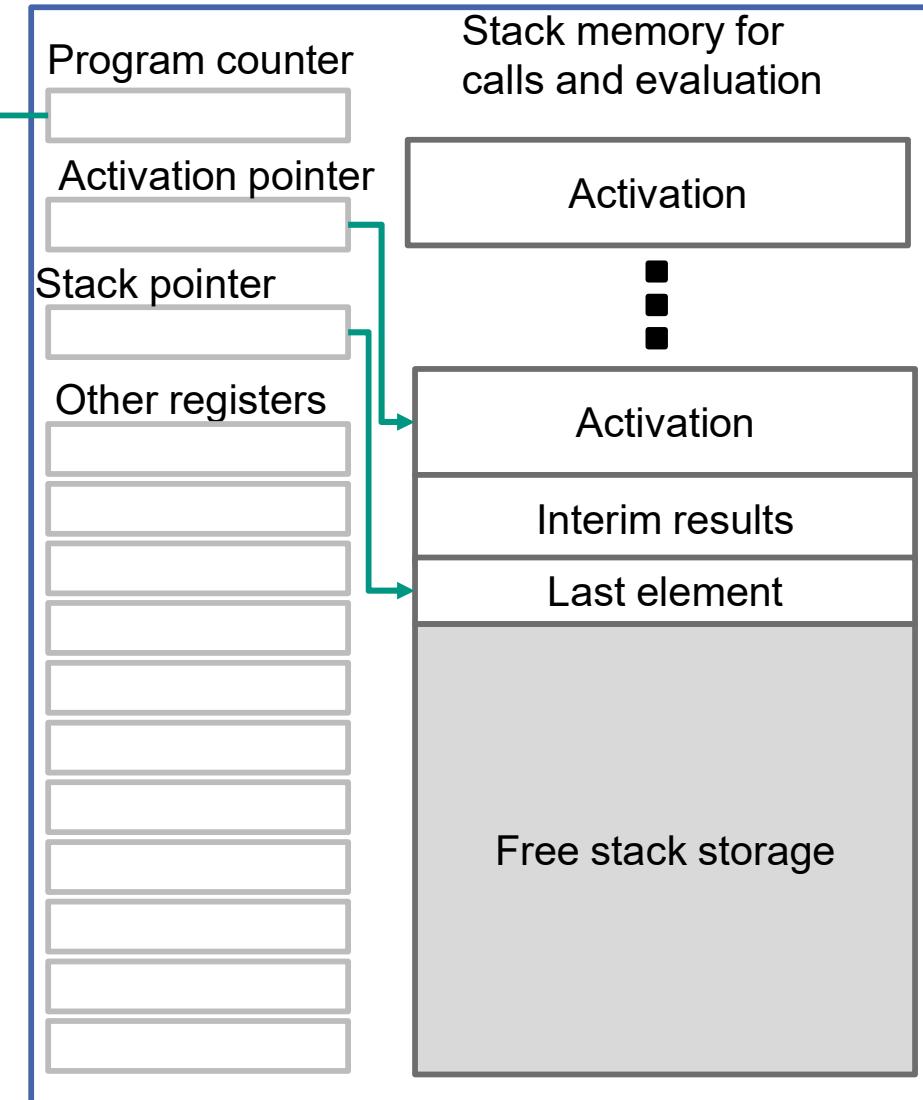
- Instruction sequence that is executed
- Exists in a process
- A thread has its own
 - Command pointer
 - Stack
 - Register
- Shares with other threads in the same process
 - Address space
 - Code/data segments
 - Other resources (e.g. open files, locks, etc.)
- CPU context change between threads of the same process faster than between different processes.

Processes and control threads

Process

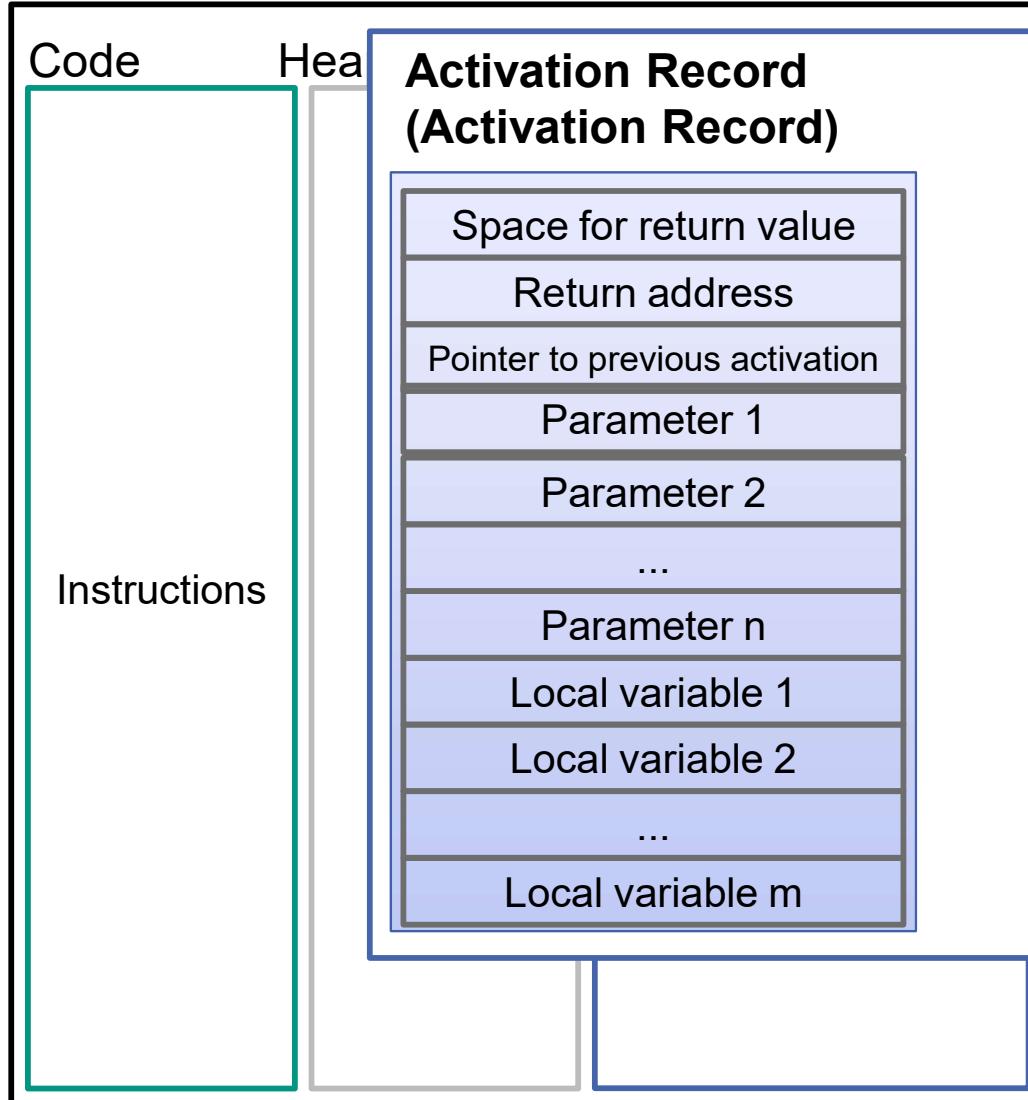


Thread 1

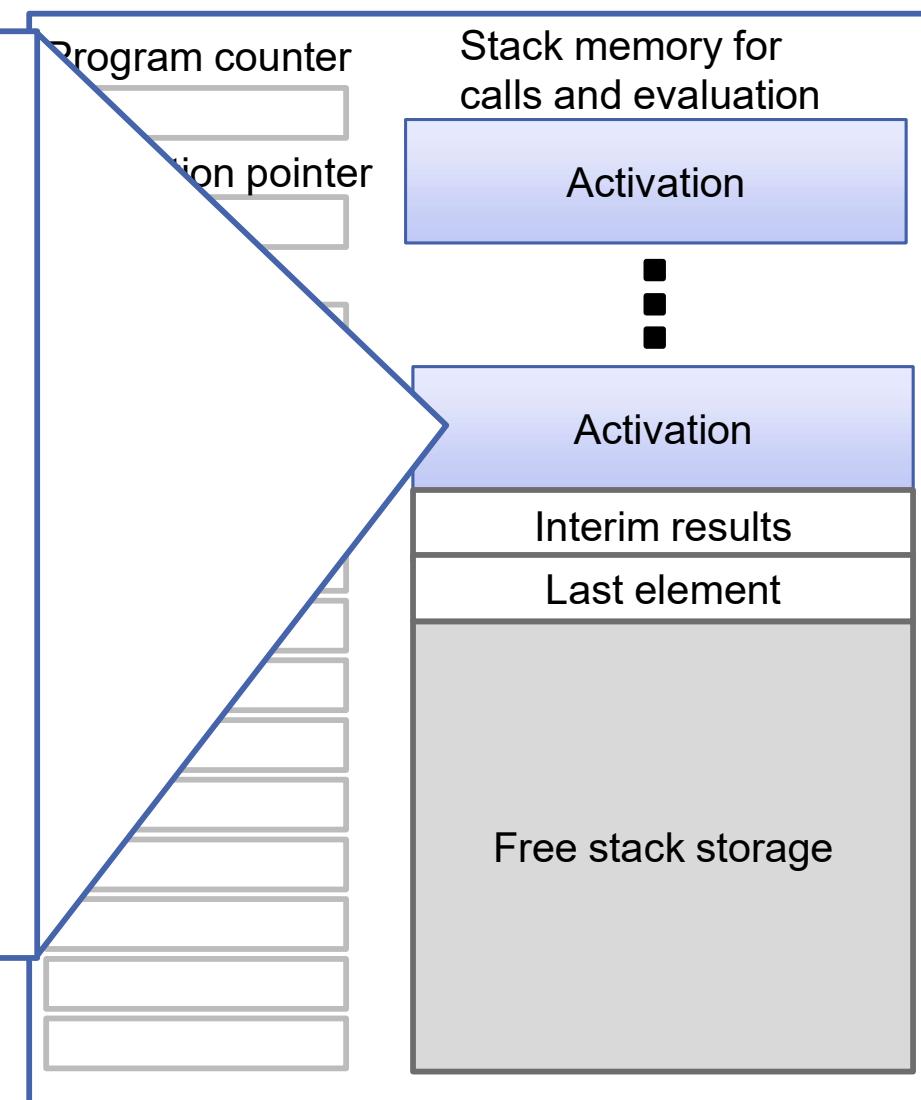


Processes and control threads

Process



Thread 1



Shared memory (2)

- Basic procedure for shared memory:
 - Threads perform parallel **tasks** (instructions)
 - Exchange of information about shared **variables** in the memory
 - In Java, this is only possible via objects on the heap. The stack is thread-local.
 - Synchronization constructs coordinate execution in the event of data or control dependencies
 - Processes and their threads are always generated by the **operating system** and distributed to processors or cores
 - Interface usually built into programming languages (e.g. Java threads) or via libraries (C Pthreads)

Flynn's Taxonomy

- {Single, Multiple} Instruction/Program x {Single, Multiple} Data
 - SISD, SIMD, MIMD, MISD
- SISD: Single Instruction x Single Data
 - von Neumann architecture:
a single instruction stream operates on a single memory
- SIMD: Single Instruction x Multiple Data
 - One instruction is applied on homogeneous data (e.g. an array)
 - e.g. vector processors of early supercomputers
- MIMD: Multiple Instruction x Multiple Data
 - Different processors operate on different data
 - e.g. today's multi-core processors
- MISD: Multiple Instruction x Single Data
 - Multiple instructions are executed simultaneously on the same data
 - e.g. in redundant architectures or in pipelines of modern processors (controversial view)

Challenges

- The programmer is responsible for partitioning the problem
- Two approaches:
 - **Task parallelism**: functional decomposition
 - Define tasks that can be executed in parallel
 - Tasks should be as independent as possible
 - **Data parallelism**: data decomposition
 - Partition the data on which the same operation is executed in parallel
 - Tasks should be able to work on the partitions as independently as possible
- „Existing“ challenges are still relevant, such as –
 - inter-process communication
 - UMA vs. NUMA ((non-)uniform memory access) and data exchange
 - Threads, shared memory, message passing
 - potentially arising race conditions and deadlocks
 - Can be avoided with techniques for mutual exclusion, such as monitors or semaphores

Theoretical Example Parallel Addition (1)

- Ideally, problems can be partitioned, such that the partitions are independent and can be processed in parallel
 - e.g. many algorithms for image processing or a vector addition

- Even a piecewise parallelization is possible:
 - Given the following array of integers

$a[] = \{4, 2, 5, 7, 3, 1, 6, 8\}$

adding the numbers sequentially requires eight cycles

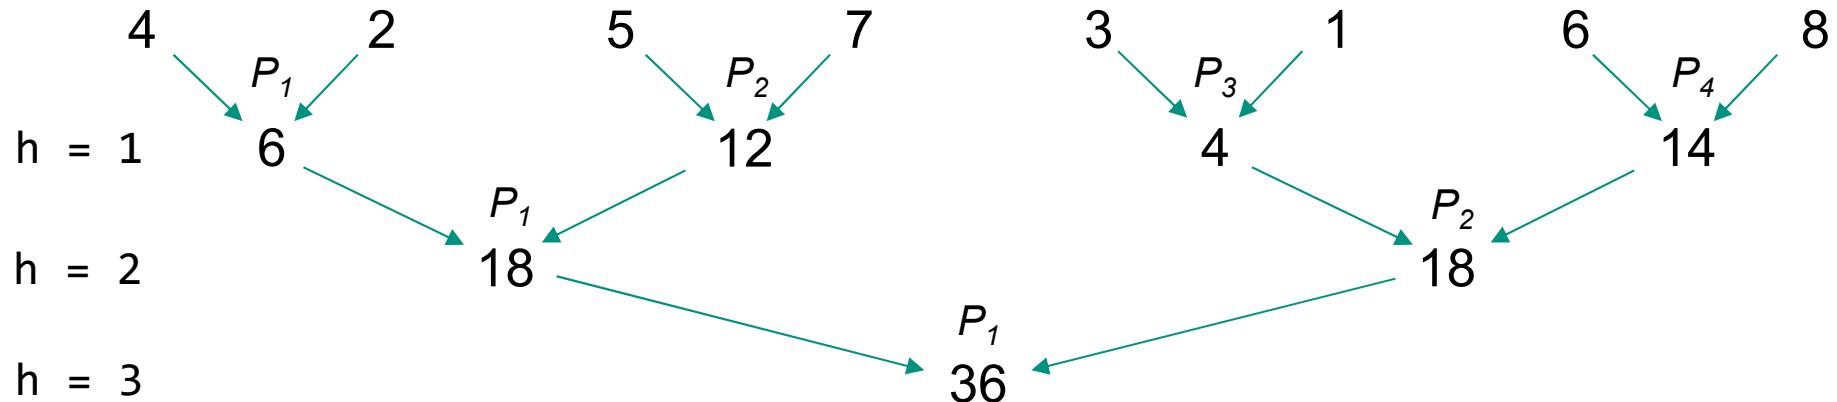
- e.g. within a for-loop
- $4 + 2 + 5 + 7 + 3 + 1 + 6 + 8 = 36$
- For n numbers, the operation requires $O(n)$ cycles

- How can we sum up “in parallel”? (e.g. on 4 processors)



Parallel Addition (2)

- Summing up the entries of $a[]$ in parallel using 4 processors
 - Per divide and conquer on a SIMD architecture



The algorithm can be executed in $O(\log n)$

```
for h := 1 to log2(n) do
    for i := 0 to n / 2h - 1 pardo
        a[i] := a[2i] + a[2i + 1]
    endfor
endfor
```

This instruction
can be distributed
on the different P_i



cf. [Wagner08]

Amdahl's Law

- The possible speedup of an algorithm executed on n processors can be computed as follows

$$S(n) = \frac{T(1)}{T(n)} = \frac{\text{execution time if processed by one processor}}{\text{execution time if processed by } n \text{ processors}}$$

- With **Amdahl's law** (1967) the **maximal** speedup that can be achieved with parallel processing on n processors can be computed
 - p = parallelizable percentage of the program

$$S(n) = \frac{1}{(1 - p) + \frac{p}{n}}$$

- Amdahl's law shows the speedup limitation due to the non-parallelizable part of a program since even for large n
- Due to synchronization and communication overhead, this speedup is hardly reachable in practice

Example: Sorting Numbers Sequentially

- ... with bubble sort
 - Given the following array of integers

$$a[] = \{7, 3, 6, 5, 8, 1, 4, 2\}$$

- An optimized sequential bubble sort requires $7 + 6 + 5 + 4 + 3 + 2 + 1 = 28$ comparisons in this example
- In general, the algorithm requires $n(n - 1)/2$ operations → $O(n^2)$



Turing Award winner
Donald Knuth

Bubble sort seems to have nothing to recommend it, except a catchy name and the fact that it leads to some interesting theoretical problems.



Serial Java Bubble Sort

```
public class SerialBubbleSort {  
    static int[] a = new int[]{7, 3, 6, 5, 8, 1, 4, 2};  
  
    public static void main(String[] args) {  
        for (int i = a.length-1; i > 0; i--) {  
            for (int j = 0; j < i; j++) {  
                if (a[j] > a[j+1]) {  
                    int temp = a[j+1];  
                    a[j+1] = a[j];  
                    a[j] = temp;  
                } // endif  
            } // for j  
        } // for i  
    } // main  
} // eof
```



Example: Sorting Numbers In Parallel

- ... with bubble sort
 - Given the following array of integers

$a[] = \{7, 3, 6, 5, 8, 1, 4, 2\}$

```
for h := 0 to n div 2 do           → O(n)
    for i := 0, 2, 4, ... pardo   → O(1)
        compareAndExchange(a[i], a[i+1])
    endfor
    for i := 1, 3, 5, ... pardo   → O(1)
        compareAndExchange(a[i], a[i+1])
    endfor
endfor
```

if (a[i] > a[i+1])
swap(a[i], a[i+1])

Each loop run is in $O(1)$, because it is completely executed in parallel on n processors



Conclusion

- Moore's Law for clock-rate is at an end → Further performance improvements primarily through parallel computing
- Two parallel processing approaches
 - For computer systems with **shared memory**
 - For computer systems with **distributed memory**
- Parallelization through **task parallelism** and **data parallelism**
- **Flynn's Taxonomy**
 - {Single, Multiple} Instruction x {Single, Multiple} Data
 - Multi-core processors allow MIMD
- **Amdahl's Law**
 - Gives the theoretical speedup of a program executed in parallel on n processors
 - Possible speedup is limited by the part of a program that cannot be parallelized

Literature and References

- [Wagner08] http://i11www.iti.uni-karlsruhe.de/_media/information/scripts/scripte/algotech08.pdf
- [Intel] <http://spmarchitecture.com/ivy-bridge-architecture/ivy-bridge-architecture-layout-147275/>
- [NASA] https://www.nasa.gov/sites/default/files/goddard_servers.jpg
- [Oracle] <https://docs.oracle.com/cd/E19455-01/806-5257/6je9h032b/index.html>