

Programmierparadigmen

Prof. Dr. Ralf Reussner

Topic 6

Design by Contract

DEPENDABILITY OF SOFTWARE-INTENSIVE SYSTEMS
INSTITUTE OF INFORMATION SECURITY AND DEPENDABILITY, FACULTY OF INFORMATICS

dsis.kastel.kit.edu



Overview on Today's Lecture

- Content
 - Motivation for design by contract
 - Formal foundation: Hoare Triples
 - Contracts for methods
 - Programming contracts with JML
 - Contract checking and tools
- Learning goals: participants –
 - understand what design by contract is
 - can argue the benefits of design by contract
 - are able to understand a contract and identify if and where it is violated
 - know about languages and tools for defining, validating and verifying contracts
 - are able to define simple contracts in JML

A Simple Stack Class

- Is this class correct?

```
class Stack {  
    private Object[] elements = new Object[10];  
    private int size = 0;  
    void push(Object element) {  
        this.elements[size++] = element;  
    }  
    public Object pop() {  
        return this.elements[--size];  
    }  
    public Object top() {  
        return this.elements[size-1];  
    }  
}
```

- Is this usage correct?

You may argue that it
is not, *but why?*

```
Stack myStack = new Stack();  
myStack.push(new Object());  
myStack.pop();  
myStack.pop();
```



Software Correctness

- What is software?
 - Artefacts created during the software development process
 - E.g. specifications, design-documents, programs
- What is *correct* software?
 - Software fulfilling its **specification**
→ Correctness is a relative notion
 - “Software verified against its specification”
 - In our previous example, there was no specification
(although you probably had one in mind)
- Validation and verification
 - Validation: Test the software with exemplary inputs (covering the most relevant equivalence classes of input data)
 - Verification: Prove software correctness for all possible input data (universal quantification)

What is the Specification of Stack?

- Informally: Manipulation of a data structure with pop and push operations according to the LIFO principle
- Additional information is necessary:
 - What if the stack is empty when calling pop or top?
 - What if the stack is full when calling push?
 - What if null is tried to be pushed?
 - What happens if several operations are called simultaneously?
 - etc.
- Informal specifications via documentation and method/parameter naming are often incomplete, inconsistent, contradictory and can only be manually checked
- A formal specification allows automated correctness checks, but is difficult to define and expensive

Specifications via Contracts

- Contracts are defined between a supplier and a client of a service about their responsibilities
- Participants:
 - The supplier of a service
 - In our example the class Stack
 - The client using a service
 - In our example the code calling operations of the class Stack
- Responsibilities:
 - Who has to do what?
 - Under which circumstances?

Formal Foundation: Hoare Triples

- Hoare triples are a central feature of the Hoare logic
 - A formal system for checking semantic correctness
 - Proposed by Sir Tony Hoare in 1969
- Form of a Hoare triple: $\{P\} C \{Q\}$
 - P: precondition
 - C: series of statements
 - Q: postcondition
- Semantics: If P is true before the execution of C, then Q is true after executing C
- Trivial example: $\{x \geq 9\} x := x + 5 \{x \geq 13\}$

Contracts for Methods

- In software, contracts can be used to specify the semantics of methods
 - **Precondition:** Specification what the supplier expects from the client
 - **Statements:** The method body
 - **Postcondition:** Specification of what the client can expect from the supplier
if the precondition is fulfilled
- Contracts do not only specify the behavior of a method but also how they have to be used by a client
 - Client has to ensure that the precondition is fulfilled
 - Client has to correctly handle the guaranteed result
- **Design by contract** idea from Bertrand Meyer [Meyer1992, Meyer1997]
 - Contractual use of methods
 - The supplier guarantees the postcondition if the precondition is fulfilled (by the client)
 - Implemented in the *Eiffel* programming language

Preconditions and Postconditions for Stack

- Preconditions:
 - push may not be called if the stack is full
 - pop / top may not be called if the stack is empty
- Postconditions:
 - After calling push, the stack may not be empty, the top element is the one that was pushed and its number of elements was increased by one
 - After calling pop, the stack may not be full and its number of elements has been decreased by one
 - After calling top, nothing has to be changed
- Are these pre- and postconditions complete?

```
public void push(Object element) {  
    size++;  
    for (int i = 0; i < size; i++) {  
        this.elements[i] = element;  
    }  
}
```

This code would also fulfill the specified postcondition



Obligations and Benefits

■ Method push example:

	Obligations	Benefits
Client	<i>Satisfy precondition:</i> Only call push(x) on a non-full stack	<i>From postcondition:</i> Get stack updated: not empty, x on top, size increased by one
Supplier	<i>Satisfy postcondition:</i> Update stack representation to have x on top, size increased by one, not empty	<i>From precondition:</i> Simpler processing thanks to the assumption that stack is not full <small>adapted from [Meyer1997]</small>

- With the precondition that the stack must not be full when calling push, the Stack class must not deal with that case
 - Alternatively, the precondition could be relaxed, but then the behavior of the stack in that case would have to be specified

Non-redundancy principle [Meyer1997]

The body of a routine shall not test for the routine's precondition.

- Often, code contains assertions or checks the precondition and throws exceptions if they are not fulfilled
- Defining contracts and checking the code against them makes those additional checks obsolete
- **Precondition design:**
 - There are two possibilities to deal with a precondition for a method:
 1. **Demanding:** Assign the responsibility to the client by adding the condition to the routine's precondition. Then the client needs to deal with cases where the pre-condition is not met.
 2. **Tolerant:** Assign the responsibility to the supplier by adding a conditional instruction and handling of violations to the control flow of the method body
 - The non-redundancy principle demands that the condition enforcement should only be assigned to either the client or the supplier



Abstraction and Intent

- Contracts abstract from concrete implementations
 - The implementation is irrelevant for a service signature and its contract
 - E.g. a square root calculation can be implemented using linear or binary search, newton's method etc.
- Different implementations can have varying non-functional properties (efficiency, memory usage etc.) but provide equal functionality
- Code itself also defines an implicit contract, but a poor one
 - No separation between contract and implementation decisions
 - Once the implementation changes, the contract might change unintendedly, too.
- Explicit contracts have several benefits
 - Allow to specify intent
 - Allow to change implementation details
 - Require a client to only understand the contract, not the implementation
 - Can be used to verify the code against them

Programming Contracts

- Several programming languages support design by contract –
 - either with pre- and postconditions as first-level entities
 - or as language extensions
- Different degrees of design by contract support:
 - Documentation: Contracts are purely informative
 - Validation: Contracts are used to generate test cases to validate if the code works as expected
 - Verification: Contracts are used to statically check if the code fulfills them
- Examples for tools and languages:
 - Eiffel: programming language with integrated contracts (by Bertrand Meyer)
 - Object Constraint Language (OCL): specialized constraint language, especially used to define contracts in the UML
 - Java Modeling Language (JML): language for specifying contracts in Java

Programming Contracts: Eiffel

- Design by contract is fully implemented in the Eiffel language
- Method pop example:

```
pop(): Object is
    require
        size > 0
    do
        ... pop logic ...
    ensure
        size = old size - 1
    end
```

References the value of
an expression before the
method was executed

- require and ensure define the pre- and postcondition
- The example contracts lacks:
 - A guarantee about the non-modification of the rest of the stack
 - A guarantee about the identity of the returned element

Java Modeling Language (JML)

- Formal behavioral interface specification language for Java
- Based on the design by contract idea
- Pre- and postconditions can be defined in specialized Java comments (all lines starting with @) in extended Java expression syntax
- Essential keywords are required for preconditions and ensures for postconditions
- Simple example for the pop method:

```
/*@ requires size > 0;
 *@ ensures  size == \old(size) - 1;
 */
Object pop() { ... pop logic ... }
```

Basic JML Syntax (1)

Syntax extensions to Java expressions:

Syntax	Meaning
$a ==> b$	a implies b
$a <==> b$	a iff b
$a <=!=> b$	$!(a <==> b)$
\result	Result of the method call
$\text{\old}(E)$	Value of E in the state before method execution

} Only available in postconditions

Extended contract for pop method:

```
/*@ requires size > 0;
 @ ensures size == \old(size) - 1 && \result == \old(top());
 @*/
Object pop() { ... pop logic ... }
```

It is possible to combine conditions with either the *and* operator or by starting a new requires / ensures lines



Basic JML Syntax (2)

- It would also be possible to specify that pop returns null if the stack is empty:

Weakest kind of precondition:
Can be omitted

```
/*@ requires true;
 @ ensures size > 0 ==> \result == \old(top());
 @ ensures size <= 0 ==> \result == null;
 @*/
Object pop() { ... pop logic ... }
```

Quantifiers in JML

- JML provides universal and existential quantifiers:
 - (`\forall declaration; range-expression; body-expression`)
 - (`\exists declaration; range-expression; body-expression`)
- We can now ensure that pop does not change other stack elements:

```
/*@ requires size > 0;
@ ensures  size == \old(size) - 1
@ ensures  \result == \old(top())
@ ensures  (\forall int i; 0 <= i && i < size;
            \old(elements[i]) == elements[i]);
@*/
Object pop() { ... pop logic ... }
```

- There are also quantifiers `\max`, `\min`, `\sum` etc.

JML: Side Effect Restrictions

- JML provides further keywords to restrict allowed side effects:
 - assignable: Specifies fields that are allowed to be written by the method
 - Per default, every field is assignable
 - The opposite of the default case can be specified as: assignable \nothing
 - pure: a method declared as pure has no side effects
 - i.e. the method does not modify any data
 - Implies assignable \nothing
- We can use that to define that calling top has no side effects:

```
/*@ requires size > 0;  
 @ assignable \nothing;  
 @*/  
Object top() { ... top logic ... }
```

We could also use @pure

- Makes postconditions on fields (especially the elements array and size) obsolete

Exceptions

- JML allows to specify expected exceptions
 - signals: defines a postcondition for the case that an exception is thrown
 - Signature: signals (E e) P
 - If an exception of type E is thrown, the predicate P must be true
 - Only exceptions defined via signals are allowed to be thrown
 - Only defines that an exception *may be* thrown, not that it *must be* thrown
 - signals_only: defines which exceptions may be thrown by the method
 - Signature: signals E1, E2, ...
 - A short form for signals without defining a condition
 - Signals implies signals_only without additional conditions
- Expecting a NullPointerException if pushing null on the stack:

```
/*@ requires size < 10;
 @ signals  (NullPointerException npEx) element == null;
 @ ensures ...
 @*/
void push(Object element) { ... push logic ... }
```

Does **not** mean that an exception is thrown if the input value is null!

Class Invariants

- Class invariants define conditions of the containing class that must hold in all user visible states
 - Must hold after each constructor execution (constructor postcondition)
 - Must hold before and after each method execution (method pre- and postcondition)
- Allow to define acceptable and consistent object states

```
class Stack {  
    //@ invariant size >= 0 && size <= 10;  
  
    private Object[] elements = new Object[10];  
    private int size = 0;  
    ...  
}
```

Class Correctness

- We can finally define when a class C is correct:
 - Let INV_C be the invariants of class C
 - Let CONSTR_C be the constructors of class C
 - Let METHODS_C be the methods of class C
 - Let PRE_M be the preconditions, BODY_M the operation body and POST_M the postconditions of a method M

Class C is correct

\Leftrightarrow

$$\begin{aligned} \forall I \in \text{CONSTR}_C : & \{ \text{PRE}_I \} \text{ BODY}_I \{ \text{POST}_I \text{ and } \text{INV}_C \} \quad \text{and} \\ \forall M \in \text{METHODS}_C : & \{ \text{PRE}_M \text{ and } \text{INV}_C \} \text{ BODY}_M \{ \text{POST}_M \text{ and } \text{INV}_C \} \end{aligned}$$

- To be precise, it would also necessary to demand that all fields of class C have default values in the precondition of the constructors

Liskov Substitution Principle

- The Liskov substitution principle restricts the possible pre- and postconditions of an overwriting method
 - Preconditions must not be more restrictive than those of the overwritten method: $\text{Precondition}_{\text{Super}} \Rightarrow \text{Precondition}_{\text{Sub}}$
 - Postconditions must be at least as restrictive as those of the overwritten methods: $\text{Postcondition}_{\text{Sub}} \Rightarrow \text{Postcondition}_{\text{Super}}$
- Regarding a complete class, the following rules apply:
 - Pre- and postcondition relations must hold for all methods as stated above
 - The class invariants must be at least as restrictive as those of the superclass: $\text{Invariants}_{\text{Sub}} \Rightarrow \text{Invariants}_{\text{Super}}$
- A special case of this is known from parameter and return types of methods (Co-/Contravariance)

Liskov Substitution Principle: Example

```
class Stack {  
  
    //@ requires element != null;  
    //@ ensures ...  
    void push(Object element) {  
        // ... push logic ...  
    }  
  
    //@ requires ...;  
    //@ ensures \result != null;  
    Object pop() {  
        // ... pop logic ...  
    }  
}
```

```
class NullAcceptingStack  
extends Stack {  
  
    //@ requires true;  
    //@ ensures ...  
    void push(Object element) {  
        // ... push logic ...  
    }  
  
    //@ requires ...;  
    //@ ensures true;  
    Object pop() {  
        // ... pop logic ...  
    }  
}
```

- Does this code fulfill the Liskov substitution principle?
No, because the postcondition of pop gets weaker in the subclass



Precondition availability rule [Meyer1997]

Every feature appearing in the precondition of a routine must be available to every client to which the routine is available.

- With this rule, every client is able to check the precondition of a routine
- This rule does not apply to postconditions
 - Postconditions can reference features that are not available to a client
 - Such postconditions are not usable by clients
 - But be careful: There is a risk to define postconditions on internal behavior, which reduces the changeability of implementation details

Precondition Availability: Example (1)

- Recall our stack implementation and our contracts for it:

```
class Stack {  
    private Object[] elements = new Object[10];  
    private int size = 0;  
  
    /*@ requires size > 0;  
     * @ ensures  size == \old(size) - 1  
     * @ ensures  \result == \old(top())  
     * @ ensures  (\forall int i; 0 <= i && i < size;  
     *             \old(elements[i]) == elements[i]);  
     */  
    Object pop() { ... pop logic ... }  
    ...  
}
```

Be aware that this is
also internal state

- Does this implementation adhere to the precondition availability rule?
No, because size is not accessible by the client, so we cannot determine whether the stack is full or empty



Precondition Availability: Example (2)

- Simple solution: add a method for retrieving the number of elements

```
class Stack {  
    private Object[] elements = new Object[10];  
    private int size = 0;  
  
    /*@ requires getElementCount() > 0;  
     * @ ensures  getElementCount() == \old(getElementCount()) - 1  
     * @ ensures  \result == \old(top())  
     * @ ensures  (\forall int i; 0 <= i && i < getElementCount();  
     *             \old(elements[i]) == elements[i]);  
     */  
    Object pop() { ... pop logic ... }  
    public int getElementCount() {  
        return size;  
    }  
  
    ...  
}
```

- Another possibility is to define `isEmpty` and `isFull` methods

Non-pure Methods in Contracts

- It is not allowed to reference methods causing side effects in contracts
 - Evaluating the condition would result in a state change of the object
→ Our newly defined method (and also the top method) must be free of side effects
 - Referenced methods must be defined as pure:

```
public /*@ pure */ int getElementCount() {  
    return size;  
}
```
- Declare methods as pure / functional as often as possible

Null-Values in JML

- JML can only interpret expressions when their assertion does not cause exceptions and yields the value „true“
 - → Dereferencing null, problematic
 - → Default Case in JML: reference types cannot be `null`
- JML introduces the `nullable` and `non_null` annotation for declarations whose type is a reference type
- Declarations are implicitly `non_null` by default unless they are annotated with the `nullable` modifier
- Possibility of `nullable` by default by annotating the outer most class or interface with `nullable_by_default`

```
/*@ requires true;
@ ensures size > 0 ==> \result == \old(top());
@ ensures size <= 0 ==> \result == null;
*/
/*@ nullable */
Object pop() { ... pop logic ... }
```

Internal Data in Contracts

- JML provides two further ways to reference private fields in contracts:
 - Specify the field as public only for the specification:

```
private /*@ spec_public */ int size = 0;
```

- Specify the condition as private behavior

```
/*@ ensures \result != null;
@ private behavior
@ requires size > 0;
*/
Object pop() { ... pop logic ... }
```

Every expression that follows this line belongs to private behavior

- This is also necessary for postconditions referencing private fields
 - Also protected behavior can be defined
- Attention:** Examples prior to this slide may not be used in an analysis due to missing elements such as private behavior or nullable-Annotations

Checking Contracts

- Contracts can be used to define a specification
- Contracts can additionally be used to check code against them if they are specified in a formal language
- Different ways of checking:
 - Perform a static contract check / verification (no code execution)
 - Perform a dynamic contract check (convert contracts into runtime checks / assertions)
 - Generate test cases from the specification

Checking JML Contracts

- For JML, several tools for checking contracts exist
 - Static contract checking, e.g. using KeY [Beckert2007] or `jml` [JML]
 - Runtime contract checking, e.g. via generated Java assertions in `jmlc` [Bhorkar2000]
 - Testing, e.g. via generated Junit test in `jmlunit` [Cheon2002]
- OpenJML (<http://www.openjml.org/>)
 - Support for static (esc) and runtime checking (rac)
 - Integrates popular logic solvers for static checking
 - Is available as command-line tool but also has an Eclipse integration and can be installed from an updatesite
 - Unfortunately, it is not capable of statically verifying our Stack example
- For details on syntax checking with OpenJML, please take a look at the appendix

Assertions as Contracts

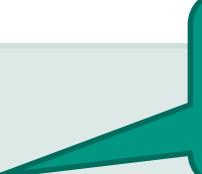
- The Java assert keyword can be used for simulating contracts

- Syntax:

```
assert expression1;
assert expression1 : expression2;
```

- If expression1 is not true, an exception containing the value of expression2 is thrown
- Method push example:

```
Object pop() {
    assert size > 0;
    int _oldSize = size;
    Object _oldTop = top();
    // ... pop logic ...
    assert size == _oldSize - 1;
    assert result == _oldTop;
    return result;
}
```



References to old values
have to be stored before
executing the method logic

Assertions as Contracts: Discussion

- Assertions can be used to check contracts at runtime
 - However, they are bad for specification purposes
 - They are restricted to Java expressions
- Assertions are not control structures
 - Conditional instructions (if-then-else) are appropriate control structures
 - Assertions only check a contract that must always hold during runtime

Assertion violation rule [Meyer1997]

A runtime assertion violation is the manifestation of a bug in the software.

- A violated precondition is the manifestation of a bug in the client
- A violated postconditions is the manifestation of a bug in the supplier, if the precondition is precise (In case of weak preconditions the bug may also be in the client and slipped through the weak precondition)
- Assertions are not an input checking mechanism
 - Assertions have to hold for software-to-software communication
 - They cannot be used to check user inputs

Contracts in Software Modeling

- Contracts are useful to define the specification of a software in early development phases, e.g. during design
- Object Constraint Language (OCL) by OMG
 - Supports the definition of contracts (specialized for models)
 - Can, for example, be used in the UML for behavior descriptions, such as method contracts in class diagrams
 - Requires the definition of a context (the method) and the pre- and postconditions after the keywords pre: and post:

```
context Stack::pop(): Object
  pre: self.size > 0
  post: result = top@pre() and size = size@pre
  body: // ... pop logic ...
```

References the expression value before method execution

Design by Contract: Discussion

- Performance impact
 - Checking contracts during execution is costly (non-performant)
 - Should not fail in a released product → only execute them in debug mode
 - Contracts are not for error handling, but only report bugs
- Degree of detail
 - If a contract is complete for a (sub)program, it provides a functional specification of that (sub)program (analogue to a program written with a functional programming language)
 - No further program logic would be required, but the program could be generated from the specification
→ shift to a more abstract (functional) programming paradigm
- Contracts are most useful when a postcondition can be given that is simpler than the code computing the result (as otherwise bugs are as likely in the postcondition as in the code)
 - Example:
(better epsilon comparison)

```
//@ ensures \result * \result == a
double sqrt(double a)
```

Design by Contract: Benefits & Limitations

■ Benefits

- Developing less error-prone software through –
 - better documentation
 - reduced error-handling code
 - clear responsibilities
 - results of static and runtime checking as a debugging aid
- Possibility to verify an implementation against a formal specification (which can be necessary for certifying software)

■ Limitation: Protocols

- Complex conditions that are as error prone as the method body itself
- Design by contract does not allow directly to define protocols, such as “routine A has to be executed before routine B”
- Protocols can be realized using status variables, requiring increased development and maintenance effort

Conclusion

- Design by contract is a principle for the contractual use of methods
 - Explicit pre- and postconditions for methods
 - Based on Hoare triples (precondition, statements, postcondition)
 - Originally implemented in the Eiffel language
 - Formally defined contracts can be used for verification purposes
- Java Modeling Language (JML)
 - Specialized Java comments with a functional software specification
 - Allow to formalize behavior of Java classes and methods
 - Tools such as OpenJML provide static or runtime checking of contracts
- Benefit: Less software bugs through –
 - Clear responsibilities
 - Better documentation
 - Automated checking and verification

Literature and References

- [Meyer1992] Bertrand Meyer. Applying "Design by Contract". Computer, 25, 40-51, 1992
- [Meyer1997] Bertrand Meyer. Object-oriented Software Construction. Prentice Hall, 1997
- [Beckert2007] Bernhard Beckert, Reiner Hähnle, Peter H. Schmitt. Verification of Object-Oriented Software. The Key Approach. Springer, 2007
- [Bhorkar2000] Abhay Bhorkar. A Run-time Assertion Checker for Java using JML. Department of Computer Science, Iowa State University, TR #00-08, 2000
- [Cheon2002] Yoonsik Cheon, Gary T. Leavens. A Simple and Practical Approach to Unit Testing: The JML and JUnit Way. In: ECOOP 2002 — Object-Oriented Programming. Lecture Notes in Computer Science, Vol. 2374. Springer, Berlin, Heidelberg, 2002.
- [JML] <http://www.eecs.ucf.edu/~leavens/JML/index.shtml>

APPENDIX

Manual (Static) Contract Checking

■ Where is this contract violated?

```
public class AddTest {  
    //@ requires b > 0;  
    //@ ensures a >= 0 ==> \result == a-b;  
    //@ ensures a < 0 ==> \result == 0;  
    public int add(int a, int b){  
        if (a < 0) return 1;  
        return a-b;  
    }  
  
    public static void main(String args[]){  
        System.out.println(new AddTest().add(2,-1));  
    }  
}
```

Method does not fulfill postcondition

Caller does not fulfill precondition $b > 0$

■ General contract checking:

- Check if all callers fulfill the preconditions
- Check if the method fulfills its postconditions (provided that the preconditions are fulfilled)



Static Contract Checking with OpenJML

- OpenJML yields two warnings in this example:

```
public class AddTest {  
    //@ requires b > 0;  
    //@ ensures a >= 0 ==> \result == a-b;  
    //@ ensures a < 0 ==> \result == 0;  
    public int add(int a, int b){  
        if (a < 0) return 1;  
        return a-b;  
    }  
  
    public static void main(String args[]){  
        System.out.println(new AddTest().add(2,-1));  
    }  
}
```

Cannot establish an assertion for this postcondition

Cannot establish an assertion: precondition is false

- To execute OpenJML with the Z3 solver [1] on the example program:

```
java -jar $OJ/openjml.jar -exec $Z3/z3.exe -esc AddTest.java
```

Path to the OpenJML folder

Path to the Z3 solver folder

Runs the static checker

[1] <http://www.openjml.org/downloads/>

Dynamic Contract Checking with OpenJML

- Executes the program with runtime checks
 - Per default, violations are reported but the program proceeds
 - Several options provide other behavior
- The dynamic contract checking requires two steps:
 1. Compile the program with the rac option of OpenJML:

```
java -jar $OJ/openjml.jar -rac AddTest.java
```

Path to the OpenJML folder

Adds additional statements for runtime checking

2. Execute the program with OpenJML added to the classpath:

```
java -cp $OJ\openjml.jar:.. AddTest
```

Path to the OpenJML folder

Replace ":" with ";" on Windows systems

- Runtime checking will only yield violations that occur in the executed code, so there is no verification of conditions